

VIETNAM NATIONAL UNIVERSITY, HO CHI MINH CITY
UNIVERSITY OF TECHNOLOGY
FACULTY OF COMPUTER SCIENCE AND ENGINEERING



Distributed System Lab

Lab 2: Introduction to Apache Spark & RDD-based Programming

Student: Nguyen Hoang Anh Thu - 2053478

HO CHI MINH CITY, NOVEMBER 2023



Contents

1	Task 1: Filtering Wrong Records	2
2	Task 2: Preprocessing	3

1 Task 1: Filtering Wrong Records

```
1 from pyspark.sql import SparkSession
2 from datetime import datetime
3 import pytz
4
5 spark_session = SparkSession.builder.appName("Log Processing").getOrCreate()
6 log_data = spark_session.sparkContext.textFile("FPT-2018-12-02.log")
7
8 # filter correct records
9 def filter_correct_records(line):
10     fields = line.split(" ")
11     criteria = len(fields) == 7 and float(fields[0]) >= 0 and fields[6].isdigit() and
12         int(fields[6]) > 0 and fields[2] != "-"
13     return criteria
14
15 # convert time format
16 def convert_time(line):
17     fields = line.split(" ")
18     input_data = fields[3] + " " + fields[4]
19     time_format = "[%d/%b/%Y:%H:%M:%S %z]"
20     try:
21         timestamp = datetime.strptime(input_data, time_format).replace(tzinfo=pytz.
22             UTC).timestamp()
23         return timestamp
24     except ValueError:
25         return None
26
27 # filter correct records
28 filtered_data = log_data.filter(filter_correct_records)
29 filtered_data = filtered_data.filter(lambda x: convert_time(x) is not None)
30 sorted_data = filtered_data.sortBy(convert_time)
31 print("Top 10 correct records:")
32 for record in sorted_data.take(10):
33     print(record)
34
35 # filter incorrect records
36 def filter_incorrect_records(line):
37     return not filter_correct_records(line)
38
39 # filter incorrect records
40 filtered_fail_data = log_data.filter(filter_incorrect_records)
41 filtered_fail_data = filtered_fail_data.filter(lambda x: convert_time(x) is not None)
42 sorted_fail_data = filtered_fail_data.sortBy(convert_time)
43 print("\nTop 10 incorrect records:")
44 for record in sorted_fail_data.take(10):
45     print(record)
46
47 # stop spark session
48 spark_session.stop()
```

Listing 1: Task 1 code

Results:

```
1 CodeCache: size=131072Kb used=12638Kb max_used=12650Kb free=118433Kb
2 bounds [0x0000000010a9e0000, 0x0000000010b650000, 0x000000001129e0000]
3 total_blobs=5218 nmethods=4391 adapters=740
```



```
4 compilation: disabled (not enough contiguous free space left)
5 Top 10 correct records:
6 0.000 58.187.29.147 HIT [02/Dec/2018:00:00:00 +0700] /live/prod_kplus_ns_hd/
  prod_kplus_ns_hd.isml/events(1541466558)/dash/prod_kplus_ns_hd-audio_vie
  =56000-49397873671168.dash 28401
7 0.055 113.23.26.76 HIT [02/Dec/2018:00:00:00 +0700] /live/prod_kplus_1_hd/
  prod_kplus_1_hd.isml/events(1541466464)/dash/prod_kplus_1_hd-video
  =2499968-926210122800.dash 1265928
8 0.000 118.69.60.62 HIT [02/Dec/2018:00:00:00 +0700] /live/prod_kplus_pm_hd/
  prod_kplus_pm_hd.isml/stb.mpd 39902
9 0.000 42.118.29.197 HIT [02/Dec/2018:00:00:00 +0700] /live/prod_kplus_ns_hd/
  prod_kplus_ns_hd.isml/prod_kplus_ns_hd.mpd 40102
10 0.000 14.231.34.1 HIT [02/Dec/2018:00:00:00 +0700] /
  cc8c96ca2e6b3aa3465a5e2d383c8e011543687445/box/_definst_/vtv3-high.m3u8 323
11 0.001 42.113.129.241 HIT [02/Dec/2018:00:00:00 +0700] /
  d053d51fe6c3935d4c25908089b7c4961543692607/ndvr/vtv3/_definst_/20181201/vtv3-high
  -20181201175215-611181.ts 742224
12 0.002 14.244.239.143 HIT [02/Dec/2018:00:00:00 +0700] /8
  d7863dc41865b32054cd8478430fb521543689560/box/_definst_/vtv2-high-2063449.ts
  760272
13 0.000 14.229.126.144 HIT [02/Dec/2018:00:00:00 +0700] /8
  c6bebd4edd78750291593f1756e861a1543686153/box/_definst_/vtv6-high.m3u8 323
14 0.000 113.22.7.224 HIT [02/Dec/2018:00:00:00 +0700] /live/prod_kplus_ns_hd/
  prod_kplus_ns_hd.isml/events(1541466558)/dash/prod_kplus_ns_hd-audio_vie
  =56000-49397873798144.dash 28840
15 0.000 113.179.83.143 HIT [02/Dec/2018:00:00:00 +0700] /98
  caef0ff9b853515fe1e3858397badf1543685819/box/_definst_/vtv6-high.m3u8 323
16
17 Top 10 incorrect records:
18 0.000 123.18.156.7 - [02/Dec/2018:00:00:00 +0700] /
  b74394d1bd14f1afb0c4be5b8c4c22481543619685/box/_definst_/vtv1-high.m3u8 166
19 0.000 113.161.6.128 - [02/Dec/2018:00:00:00 +0700] /4176
  f0256a9c8b9d5cf6854e27ade7461543681587/box/_definst_/vtv3-high.m3u8 166
20 0.000 42.115.220.10 HIT [02/Dec/2018:00:00:00 +0700] /live/prod_kplus_1_hd/
  prod_kplus_1_hd.isml/stb.mpd 0
21 0.000 27.3.65.112 - [02/Dec/2018:00:00:00 +0700] /1
  c68cc3b620f13797a8dc7d60ea8bec41543571805/box/_definst_/vtv3-high.m3u8 166
22 0.000 113.162.4.228 - [02/Dec/2018:00:00:00 +0700] /73
  f2e7ff5a352dc48bc3f46ccfe155a701543617776/box/_definst_/vtv3-high.m3u8 166
23 0.000 14.181.24.168 - [02/Dec/2018:00:00:00 +0700] /9
  f9acfb050af241723b7576ac79879d81543634619/box/_definst_/vtv3-high.m3u8 166
24 0.000 14.228.49.147 - [02/Dec/2018:00:00:00 +0700] /
  ab65523b72105c0ad027936fddc3c5ec1543562106/box/_definst_/vtv3-high.m3u8 166
25 0.000 123.17.19.243 - [02/Dec/2018:00:00:00 +0700] /45
  cb0917fe18ec62f2c5af70a9e992b61543632020/box/_definst_/vtv3-high.m3u8 166
26 0.000 14.189.182.247 - [02/Dec/2018:00:00:00 +0700] /
  e987966426fe66700b9a0775a08f7c2d1543548337/box/_definst_/vtv3-high.m3u8 166
27 0.000 14.165.116.89 - [02/Dec/2018:00:00:00 +0700] /6
  ab3ab2fcc3cb07a5a9f6cdb7784da9a1543675624/box/_definst_/vtv1-high.m3u8 166
```

Listing 2: Task 1 result

2 Task 2: Preprocessing

```
1 from pyspark import SparkConf, SparkContext
2 from pyspark.sql import SparkSession
```



```
3 from pyspark.mllib.stat import Statistics
4 from pyspark.mllib.linalg import Vectors
5 import numpy as np
6
7 spark_session = SparkSession.builder.appName("LogAnalysis").getOrCreate()
8 spark_context = spark_session.sparkContext
9 log_data = spark_context.textFile("FPT-2018-12-02.log")
10
11 def is_valid_record(line):
12     fields = line.split(" ")
13     criteria = len(fields) == 7 and float(fields[0]) >= 0 and fields[6].isdigit() and
14         int(fields[6]) > 0 and fields[2] != "-"
15     return criteria
16
17 valid_log_data = log_data.filter(is_valid_record)
18
19 # classify service
20 def get_service_type(line):
21     content_name = line.split(" ")[5]
22     if content_name.endswith(".mpd") or content_name.endswith(".m3u8"):
23         return "HLS"
24     elif content_name.endswith(".dash") or content_name.endswith(".ts"):
25         return "MPEG-DASH"
26     else:
27         return "Web Service"
28
29 classified_log_data = valid_log_data.map(lambda line: (get_service_type(line), 1))
30 service_counts = classified_log_data.reduceByKey(lambda a, b: a + b)
31
32 for service, count in service_counts.collect():
33     print(f"{service}: {count} records")
34
35 # extract IP
36 def extract_ip(line):
37     ip = line.split(" ")[1]
38     return ip
39
40 unique_ips = valid_log_data.map(extract_ip).distinct()
41
42 # load IP information
43 ip_info_data = spark_context.textFile("IPDict.csv")
44 ip_info_dict = ip_info_data.map(lambda line: (line.split(",")[0], (line.split(",")
45     [1], line.split(",")[2], line.split(",")[3]))).collectAsMap()
46 ip_info_broadcast = spark_context.broadcast(ip_info_dict)
47
48 # enrich log record
49 def enrich_record(line):
50     fields = line.split(" ")
51     ip = fields[1]
52     additional_info = ip_info_broadcast.value.get(ip, ("Unknown", "Unknown", "Unknown"))
53     latency = float(fields[0])
54     city = additional_info[1]
55     content_size = int(fields[len(fields) - 1])
56     return (ip, additional_info, city, latency, fields[5], content_size)
57
58 enriched_log_data = valid_log_data.map(enrich_record)
```



```
58 # unique ISPs
59 unique_isps = enriched_log_data.map(lambda log: log[1][2]).distinct().collect()
60 print(f"Number of unique ISPs: {len(unique_isps)}")
61
62 # records from Ho Chi Minh City
63 hcm_records = enriched_log_data.filter(lambda log: log[2] == "Ho Chi Minh City")
64 print(f"Number of records from Ho Chi Minh City: {hcm_records.count()}")
65
66 # traffic from Hanoi
67 hanoi_traffic = enriched_log_data.filter(lambda log: log[2] == "Hanoi").map(lambda
    log: log[5]).reduce(lambda a, b: a + b)
68 print(f"Total traffic from Hanoi: {hanoi_traffic}")
69
70 # latency statistics
71 latencies = enriched_log_data.map(lambda log: log[3])
72 latencies_vector = latencies.map(lambda latency: Vectors.dense(latency))
73 latency_stats = Statistics.colStats(latencies_vector)
74
75 print(f"Mean Latency: {latency_stats.mean()[0]}")
76 print(f"Maximum Latency: {latency_stats.max()[0]}")
77 print(f"Minimum Latency: {latency_stats.min()[0]}")
78
79 # Additional logic for median, maximum, and second maximum latency
80 sorted_latencies = latencies.sortBy(lambda latency: latency).collect()
81
82 median_index = int((len(sorted_latencies) + 1) / 2) - 1
83 median_latency = sorted_latencies[median_index]
84 print(f"Median Latency: {median_latency}")
85
86 max_latency_sorted_list = sorted_latencies[-1]
87 print(f"Maximum Latency (sorted list): {max_latency_sorted_list}")
88
89 second_max_latency = sorted_latencies[-2]
90 print(f"Second Maximum Latency: {second_max_latency}")
91
92 spark_session.stop()
```

Listing 3: Task 2 code

Results:

```
1 CodeCache: size=131072Kb used=16753Kb max_used=16768Kb free=114318Kb
2 bounds [0x00000001071e0000, 0x0000000108260000, 0x000000010f1e0000]
3 total_blobs=7137 nmethods=6236 adapters=814
4 compilation: disabled (not enough contiguous free space left)
5 MPEG-DASH: 826313 records
6 Web Service: 13976 records
7 HLS: 462938 records
8 Number of unique ISPs: 125
9 Number of records from Ho Chi Minh City: 217212
10 Total traffic from Hanoi: 204245300091
11 Mean Latency: 0.1516318983569242
12 Maximum Latency: 199.658
13 Minimum Latency: 0.0
14 Median Latency: 0.0
15 Maximum Latency (sorted list): 199.658
16 Second Maximum Latency: 119.467
```

Listing 4: Task 2 result



Link GitHub for Lab 2: [DS_Lab2_2053478](#)